PMx Weekly Digest: February 17, 2026

Paper 1: The Regularized Horseshoe for Covariate Selection Improves Convenience and Predi

Bibliographic Information

- **Title**: The Regularized Horseshoe for Covariate Selection Improves Convenience and Predictive Performance in Population PK/PD Models
- **Authors**: Pourzanjani A, Davis C
- **Journal**: CPT: Pharmacometrics & Systems Pharmacology, 2026
- **DOI**: https://ascpt.onlinelibrary.wiley.com/doi/10.1002/psp4.70198?af=R

Quick Take

A compelling case for abandoning stepwise covariate modeling (SCM) in favor of Bayesian shrinkage. The Regularized Horseshoe (RHS) prior offers a 'one-and-done' model fit that avoids the selection bias and inflated effect sizes common in standard pharmacometric workflows.

Executive Summary

This study introduces the Regularized Horseshoe (RHS) prior as a superior alternative to stepwise covariate methods (SCM, COSSAC) and LASSO for population PK/PD modeling. Addressing the pervasive issue of selection bias—where stepwise methods systematically overestimate effect sizes—the authors demonstrate that RHS simultaneously identifies relevant covariates and regularizes their coefficients in a single Bayesian model fit. Through extensive simulations and four real-world datasets (including Remifentanil and Warfarin), the RHS demonstrated superior predictive performance (LPPD) on held-out data compared to traditional methods. The authors provide a set of generalizable hyperparameters and Stan/Torsten code, making this advanced Bayesian technique accessible for routine application to prevent overfitting and improve model generalizability.

⚡ Methodological Snapshot

The authors implemented the RHS prior within a Bayesian framework using Stan/Torsten. They compared RHS against SCM, COSSAC (run in Monolix/NONMEM), and LASSO (PsN). The evaluation utilized a comprehensive simulation study varying sample size (N=12 to 300) and sparsity levels, alongside four real-world PK/PD datasets. Performance was assessed using Log Pointwise Predictive Density (LPPD) on held-out data (LOGO-CV), calibration of prediction intervals, and specific error metrics (Type 1/2, Type S for sign, and Type M for magnitude).

🏗 Structural Model Breakdown

The core innovation is the hierarchical prior structure on the covariate coefficients θ_j. The model assumes $\theta_j \sim \text{Normal}\left(0, \tau\tilde{\lambda}_j\right)$, where τ is a global shrinkage parameter (pulling all towards zero) and $\tilde{\lambda}_j$ is a local parameter allowing specific signals to escape shrinkage. The 'slab' is defined by an Inverse-Gamma distribution on the auxiliary parameter c, creating a Student-t like tail that

prevents over-shrinkage of large effects. The structural PK models varied by dataset (1-cmt and 2-cmt linear models).

Detailed Methodological Analysis

Modeling Approach

Bayesian hierarchical modeling using the Regularized Horseshoe (RHS) prior for covariate coefficients. Implemented in Stan/Torsten. Benchmarked against SCM/COSSAC (SAEM/FOCE) and LASSO.

Data Sources

Simulated 2-compartment PK data with 10 covariates (based on NHANES structure) across 12 scenarios (varying N and sparsity). Real-world datasets: Remifentanil PK (N=65), Theophylline PK (N=12), Tobramycin PK (N=97), Warfarin PK/PD (N=32).

Estimation Methods

Bayesian MCMC (NUTS sampler) for RHS (4 chains, 500 warmup, 1000 sampling). SAEM (Monolix) for simulated SCM/COSSAC. FOCE (NONMEM) for real-world SCM/COSSAC/LASSO.

Model Evaluation

Leave-One-Group-Out Cross-Validation (LOGO-CV) calculating LPPD. Bias, MSE, Type 1/2 error, Type S (sign) error, and Type M (magnitude) error.

Covariate Analysis

Simultaneous selection and estimation using RHS. Continuous covariates were standardized; positive covariates log-transformed. Comparison of 'selected' covariates based on 95% posterior intervals vs. stepwise inclusion criteria.

Scientific Context & Motivation

Covariate selection is a critical step in pharmacometrics to explain inter-individual variability. The field has long relied on stepwise hypothesis testing (SCM), which is known to suffer from selection bias, multiple testing issues, and 'Type M' (magnitude) errors where effect sizes are exaggerated. While LASSO offers regularization, it can bias large effects downward. The RHS, a continuous 'spike-and-slab' Bayesian prior, theoretically resolves these issues by shrinking noise to zero ('spike') while leaving true signals largely unbiased ('slab'). This paper bridges the gap between advanced Bayesian statistics and practical pharmacometrics.

Statistical Rigor Assessment

Exceptional. The study moves beyond simple parameter recovery to assess predictive performance (LPPD) and specific statistical pathologies like Type M error. The use of 100 replicates for simulations and rigorous cross-validation on real data provides high confidence in the results. The choice of hyperparameters was validated using prior predictive checks.

📊 Key Findings

1. **Elimination of Exaggeration:** Stepwise methods consistently overestimated effect sizes (high Type M error), particularly in small datasets (N=12), whereas RHS estimates were conservative and closer to truth.
2. **Predictive Superiority:** RHS achieved higher LPPD on held-out data in >70% of simulated replicates and across all four real-world datasets compared to SCM/COSSAC.
3. **Small Sample Robustness:** In the N=12 Theophylline dataset, SCM found no covariates, while RHS identified signal and improved prediction, demonstrating the benefit of 'soft' probabilistic inclusion over binary selection.
4. **Efficiency:** RHS requires only a single model fit, whereas SCM/COSSAC require many iterative runs.

💡 Clinical & Regulatory Implications

Adopting RHS can lead to more robust dosing algorithms by filtering out spurious covariate relationships that SCM might include due to chance (p-hacking). It allows for safer inference in small populations (e.g., pediatrics, rare diseases) where traditional methods are prone to overfitting. The resulting models are better calibrated for predicting future patient responses.

Strengths & Limitations

Strengths

- Addresses the fundamental statistical flaw of selection bias in stepwise methods.
- Provides a 'drop-in' set of hyperparameters suitable for general PK/PD models.
- Demonstrates superior predictive performance on strictly held-out data.
- Code (Stan/Torsten) is provided, lowering the barrier to entry.

Limitations (Acknowledged by Authors)

- RHS does not produce a strictly sparse model (coefficients are never exactly zero), which may require post-processing if a simplified list is needed.
- Correlated covariates may split the explanatory power in the posterior, making interpretation slightly more complex than 'selecting one'.
- Hyperparameters might need tuning for very specific non-standard problems.

Limitations (Expert Review)

- Requires transition to Bayesian framework (Stan), which has a steeper learning curve than NONMEM/Monolix for many pharmacometricians.
- Computation time for MCMC is generally longer than FOCE/SAEM, though the authors note RHS is only *one* run vs. many for SCM.
- The 'selection' criteria for RHS (95% CI excluding zero) is a heuristic comparison; the method is inherently continuous.

Generalizability
High. The hyperparameters were tuned to match sparsity patterns seen in 17 published datasets, suggesting broad applicability to standard PopPK analyses.

Key Equations
RHS Hierarchical Prior

$$\theta_j \mid \tau, \lambda_j, c \sim \text{Normal}\left(0, \tau\tilde{\lambda}_j\right)$$

The conditional distribution of the covariate coefficient, governed by global shrinkage (tau) and local shrinkage (lambda).

Effective Local Shrinkage

$$\tilde{\lambda}_j^2 = \frac{c^2\lambda_j^2}{c^2 + \tau^2\lambda_j^2}$$

Defines the effective variance for the coefficient, balancing the local signal against the regularization of the slab (c).

Shrinkage Factor

$$\hat{\theta} = (1-\kappa)\hat{\theta}_{\text{ML}}$$

Conceptual representation of how the estimate is shrunk towards zero relative to the Maximum Likelihood estimate.

Log Pointwise Predictive Density

$$\text{lppd}_{\text{logo-cv}} = \sum_{i=1}^{n} \log\left(\frac{1}{S}\sum_{s=1}^{S} p(y_i \mid \eta_{is})\right)$$

The metric used to evaluate predictive performance on held-out data (Leave-One-Group-Out Cross-Validation).

Figures & Tables

- **Table 1**: Comparison of FFEM, FREM, Stepwise, LASSO, and RHS methods.
  - *Significance*: Succinctly highlights that RHS is the only method that provides unbiased CIs, low variance, sparse estimates, and requires only a single fit.
- **Figure 3**: Posterior distributions of V1-Age effect vs. point estimates.
  - *Significance*: Visually demonstrates how RHS captures uncertainty (bimodal posterior) where SCM forces a binary decision, often leading to incorrect inclusions or exclusions.
- **Figure 5**: Evaluation metrics (Bias, MSE, Type M error) vs. Sparsity and Sample Size.
  - *Significance*: Quantifies the 'exaggeration' (Type M error) of SCM, showing it inflates effects by up to 1.4x, while RHS remains conservative.

- **Table 2**: Real-world data comparison of selected covariates and LPPD.
  - *Significance*: Shows that RHS consistently achieves better predictive density (higher LPPD) than SCM/COSSAC across diverse real datasets.

Code & Reproducibility Assessment

Excellent. The authors provide Stan/Torsten code, R scripts, and specific hyperparameter values (Data S1, S2, S4).

Supplementary Materials

Includes Stan model code (Data S1), R scripts for execution (Data S2, S4), and detailed simulation setup/results (Data S3). The availability of the Torsten implementation is a significant resource.

Future Directions

Investigation of 'Projective Predictive Inference' to reduce the full RHS posterior into a sparse subset of covariates without losing predictive power. Application of RHS to QSP models where parameter identifiability is a major challenge.

Expert Commentary

This paper marks a necessary maturation in pharmacometrics statistics. We have long known that SCM is statistically sub-optimal (effectively p-hacking), but it remained the default due to convenience. The RHS offers a mathematically sound alternative that fits naturally into the growing Bayesian trend in the field. The demonstration of reduced Type M error is particularly damning for SCM in small-N scenarios. The main hurdle will be the software shift; however, with tools like Torsten and bbr, this is becoming easier. This method should become the gold standard for covariate analysis in high-stakes or low-data environments.

Bottom Line

The Regularized Horseshoe (RHS) is a robust, one-step Bayesian alternative to Stepwise Covariate Modeling that provides superior predictive performance and avoids the inflation of covariate effect sizes. For pharmacometricians comfortable with or willing to learn Stan, RHS should replace SCM as the default method for covariate selection, especially in datasets with limited sample sizes or high dimensionality.

📊 Figures

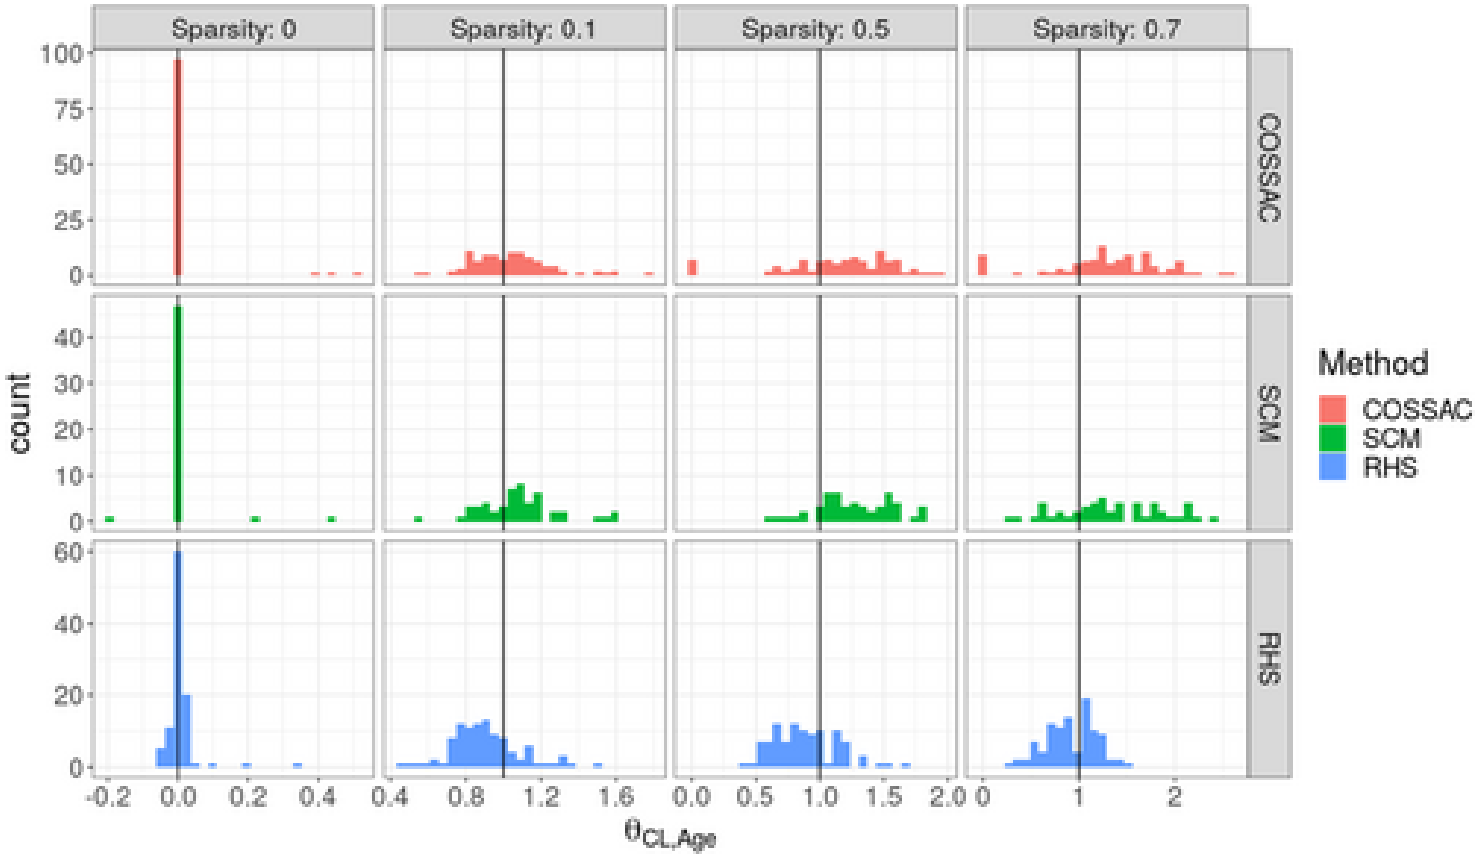

Figure 1: Figure 1

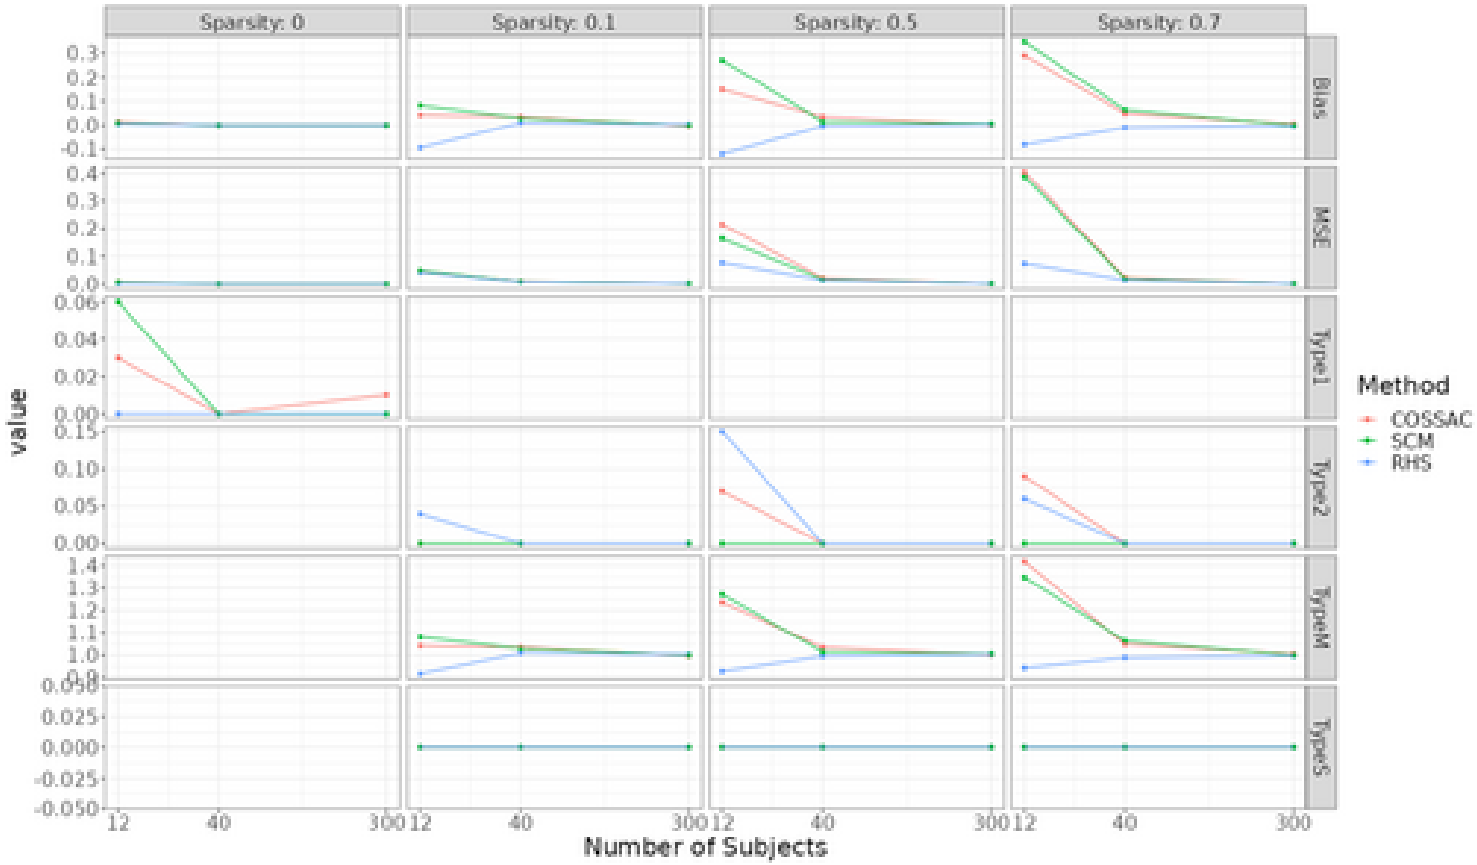

Figure 2: FIGURE 5: Evaluation metrics of methods for theta_mathrmCL,mathrmAge as a function of ground-truth sparsity value and number of su

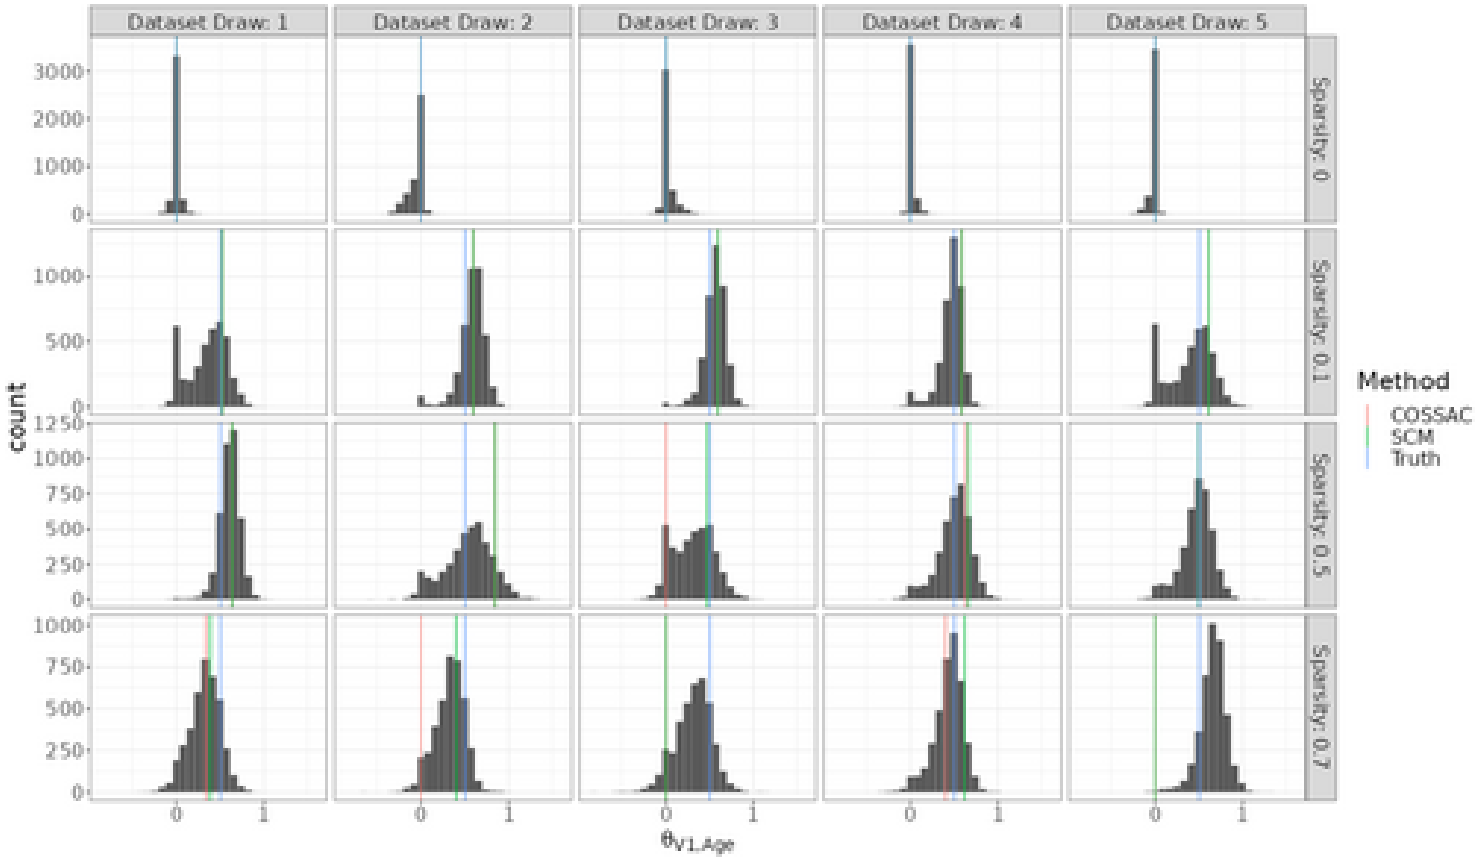

Figure 3: Details are in the caption following the image