

PMx Weekly Digest: July 13, 2026

Content Source: Full Text

Paper 1: The Empirical Bayes Variational Autoencoder—A Neural ODE Approach for Population

Bibliographic Information

- **Title:** The Empirical Bayes Variational Autoencoder—A Neural ODE Approach for Population Modeling in Pharmacology
- **Authors:** Baaz M, Sjöberg A, Jirstrand M
- **Journal:** CPT: Pharmacometrics & Systems Pharmacology, 2025
- **DOI:** 10.1002/psp4.70000

Paper Type: AI_ML

Quick Take

This paper introduces an empirical Bayes variational autoencoder with a neural ODE decoder for population PK/PD modeling. The method learns flexible latent representations of inter-individual variability and covariate effects, outperforming a fixed-prior VAE baseline in simulation studies and achieving comparable performance to a neural ODE approach on a small clinical dataset. The work is a proof-of-concept for combining variational inference with neural ODEs in pharmacometrics.

Executive Summary

This paper presents an empirical Bayes variational autoencoder (VAE) combined with a neural ODE decoder for population PK/PD modeling. The framework learns flexible latent representations of inter-individual variability, incorporates covariate-dependent population priors, and uses a transformer-based encoder for posterior inference. In simulation studies, the empirical Bayes VAE accurately recovered correlated latent structures and covariate effects, outperforming a fixed-prior VAE baseline. On a small clinical theophylline dataset, it achieved predictive performance comparable to a previously reported neural ODE approach. The authors propose an input-response normalization strategy to disentangle external inputs from latent effects. The work is positioned as a methodological proof-of-concept, highlighting both the potential and current limitations of variational neural approaches in pharmacometrics.

Scientific Context & Motivation

Classical NLME models rely on rigid parametric assumptions about structural models, variability distributions, and error models, requiring substantial manual specification. Recent ML advances (neural ODEs, VAEs) offer more flexible alternatives but face challenges in scalability, identifiability, and integration of population structure. This paper addresses the gap by combining empirical Bayes VAEs with neural ODE decoders, providing a principled way to learn correlated latent structures and covariate effects while maintaining probabilistic inference.

✂ Methodological Snapshot

The model uses a transformer-based encoder to map individual time-series data (including dosing and covariates) to a posterior distribution over latent variables. The decoder consists of a neural ODE that integrates a latent dynamical state forward in time, followed by a mapping to the concentration domain with a Gaussian additive-proportional error model. The population prior is a multivariate Gaussian with mean and covariance parameterized by a neural network conditioned on baseline covariates (empirical Bayes). Training optimizes the evidence lower bound (ELBO) via gradient descent, with gradients through the ODE solver obtained by direct backpropagation (adjoint method not needed due to short trajectories). A pretraining phase (input-response normalization) first trains the decoder with latent effects fixed to the population mean, then normalizes observations by the decoder’s prediction to reduce identifiability issues.

🧩 Task Framing

The task is population PK/PD modeling: learning a generative model that captures inter-individual variability in drug concentration-time profiles, incorporating baseline covariates and dosing schedules. This is framed as a variational inference problem where the goal is to infer a posterior distribution over individual-specific latent parameters and learn population-level priors.

Dataset Description

Simulation study: synthetic PK data from a one-compartment model with first-order absorption, repeated dosing, IIV on initial concentration, absorption rate, and elimination rate. Two scenarios: (1) uncorrelated random effects, (2) correlated random effects with a baseline covariate effect on elimination rate. Large number of individuals. Clinical case study: theophylline PK dataset from Monolix (single oral dose, ~24h sampling, small sample). 100 random 70/30 cross-validation splits.

Model Architecture

Encoder: transformer-based network. Decoder: neural ODE (fixed-step Runge-Kutta solver) with ~136,000-500,000 parameters. Latent space dimensionality: 3 (simulation) and 2 (clinical). Population prior: multivariate Gaussian with mean and covariance learned via neural network (empirical Bayes). Observation model: Gaussian with additive and proportional components.

Training Details

Optimization: gradient descent on ELBO. Gradients through neural ODE obtained by direct backpropagation through fixed-step Runge-Kutta solver (adjoint method not needed). Reconstruction term corresponds to weighted MSE with learned noise parameters. Pretraining phase: decoder trained with latent effects fixed to population mean (L2 loss). Input-response normalization: observations divided by decoder prediction under population mean. Latent dimensions: 3 (simulation), 2 (clinical). Decoder parameters: ~136,000-500,000. Computation: not specified in detail.

Evaluation Metrics

Simulation: visual comparison of population distributions, recovery of true parameters via posterior means, Random Forest prediction of individual parameters, error standard deviation estimates. Clinical: median absolute prediction error (0.42 vs 0.57 for Bräm model), trimmed MSE (0.40 vs 0.46), trimmed R^2 (0.88 vs 0.87). Model diagnostics: individual weighted residuals, observed vs predicted plots, visual predictive checks.

Deployment Caveats

Generalizability is limited by the small clinical dataset and the controlled simulation conditions. Extrapolation beyond training dosing schedules showed stable behavior only within a limited range. The method may not recover asymptotic behavior (e.g., Michaelis-Menten saturation) in sparsely observed regions. Identifiability between external inputs and latent effects requires the proposed input-response normalization. The model's performance on larger, more heterogeneous real-world datasets remains to be demonstrated.

☒ Key Findings

- 1) The empirical Bayes VAE accurately captures correlated latent structures and covariate effects in simulated PK data, while a fixed-prior VAE exhibits systematic biases.
 - 2) Input-response normalization improves extrapolation to unseen dosing schedules by disentangling external inputs from latent effects.
 - 3) On a small clinical theophylline dataset, the empirical Bayes VAE achieves predictive performance comparable to a previously reported neural ODE approach (median absolute prediction error 0.42 vs 0.57).
 - 4) The model provides probabilistic latent representations that quantify uncertainty and support prediction in new populations.
-

Strengths & Limitations

Strengths

- Clear conceptual connection between VAE and classical NLME (SAEM) framework.
- Empirical Bayes prior effectively captures correlated latent structures.
- Input-response normalization addresses a key identifiability issue.
- Probabilistic latent representation provides uncertainty quantification.
- Scalable gradient-based optimization compared to MCMC-based SAEM.

Limitations (Acknowledged by Authors)

- Small clinical dataset limits conclusions about comparative performance.
- Extrapolation far beyond observed data is limited (no mechanistic saturation effects).
- Input-response normalization is a heuristic; detailed exploration of post-training step is beyond scope.
- Model is a proof-of-concept; current limitations exist.

Limitations (Expert Review)

- No comparison with classical NLME models (e.g., NONMEM, Monolix) on the clinical dataset.

- Computational cost of training large neural ODEs is not reported.
- The claim of ‘data modality agnosticism’ is not demonstrated.
- The post-training normalization step for individual inference is not fully developed or tested.

Generalizability

Demonstrated on simulated data and one small clinical dataset; generalizability to larger, more heterogeneous real-world datasets is unproven.

Figures & Tables

- **Figure 1:** Schematic of the empirical Bayes VAE architecture, showing encoder, decoder (neural ODE), and covariate-conditioned prior.
 - *Significance:* Provides a clear visual overview of the model components and their connections.
 - **Figure 2:** Comparison of population simulations from fixed-prior VAE (top) and empirical Bayes VAE (bottom) against observed data with correlated individual parameters.
 - *Significance:* Demonstrates the key advantage of the empirical Bayes prior: capturing correlated latent structure that the fixed-prior VAE misses.
 - **Figure 3:** Model predictions for dosing schedules not seen during training (different doses and intervals).
 - *Significance:* Shows extrapolation capability within the empirical dose-time domain, with minor discrepancies only at high doses.
 - **Figure 4:** Scatter plots of true vs. inferred parameters, including Random Forest predictions.
 - *Significance:* Validates that the encoder recovers meaningful latent representations that correlate with ground-truth parameters.
 - **Figure 5:** Simulations vs. test data stratified by treatment arm and covariate, with 95% CI.
 - *Significance:* Demonstrates accurate recovery of covariate effects and correlation structure, with error estimates close to true values.
 - **Figure 6:** VPC, IWRES histogram, and observed vs. predicted plot for the theophylline case study.
 - *Significance:* Shows satisfactory model fit on clinical data, supporting practical feasibility.
-

Code & Reproducibility Assessment

Code is publicly available at <https://github.com/FraunhoferChalmersCentre/neuralODEpopulation>. The paper provides detailed architecture and training procedures in DATA S1, supporting reproducibility.

Supplementary Materials

DATA S1 provides detailed model architecture, ODE solver configuration, encoder and decoder networks, and full training procedure. DATA S2 contains additional plots for the fixed-prior VAE without input-response normalization.

Future Directions

Future work should validate the framework on larger, more diverse clinical datasets with multiple covariates and sparse sampling. Exploring hybrid models that combine mechanistic PD components with neural network hazard functions for time-to-event modeling is a promising direction. The post-training normalization step for individual-level inference needs further development and testing. Systematic comparison with classical NLME models on a range of benchmarks would strengthen the evidence base.

Expert Commentary

This paper makes a valuable contribution by explicitly connecting the VAE framework to classical NLME concepts (SAEM, random effects, empirical Bayes). The input-response normalization is a clever solution to an identifiability issue that often plagues flexible encoder-decoder architectures. However, the clinical validation is limited to a small, simple dataset; the real test will be on larger, more complex datasets with multiple covariates and sparse sampling. The claim that the model is ‘agnostic to data modality’ is intriguing but not demonstrated. The computational cost of training large neural ODEs may still be prohibitive for routine use in drug development.

Bottom Line

The empirical Bayes VAE with neural ODE decoder offers a scalable, flexible alternative to classical NLME models, capturing correlated latent structures and covariate effects without rigid parametric assumptions. Practitioners should view this as a proof-of-concept: the method shows promise for automated model building and handling complex dynamics, but requires larger datasets and careful normalization to ensure identifiability.

☒ Figures




Figure 1: Overview of the empirical Bayes variational autoencoder. The encoder learns the posterior of the individual parameters, which gives individual time course data a




Figure 2: The top row shows population model simulations (blue) from the fixed-prior variational autoencoder, based on sampling from the prior, compared to observed data (




Figure 3: Model predictions (blue) evaluation using the empirical Bayes variational autoencoder trained on the 400-mg and 800-mg dosing schedules (days 0, 8, 16, and 24),




Figure 4: The first three columns display the true model parameters alongside the prior-normalized mean posterior estimates for each individual, with treatment arms distin




Figure 5: Empirical Bayes variational autoencoder simulations compared to test data stratified on both treatment arm and covariate. The 95% confidence interval, based on 1




Figure 6: Visual predictive check, histogram of individual weighted residuals (IWRES), and observed concentrations plotted against model predictions from an empirical Bayes